

A (Very Very) Brief Introduction to Language Models

Aniello De Santo

he/him

`aniellodesanto.github.io`

`aniello.desanto@utah.edu`

CUNY GC

March 27, 2023

A Definition

Language Models assign probabilities to sequences of words.

A Definition

Language Models assign probabilities to sequences of words.

$$\begin{aligned}P(X_1 \dots X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_{1:2}) \dots P(X_n|X_{1:n-1}) \\&= \prod_{k=1}^n P(X_k|X_{1:k-1})\end{aligned}$$

$$P(w_n|w_{n-N+1:n-1}) = \frac{C(w_{n-N+1:n-1} w_n)}{C(w_{n-N+1:n-1})}$$

$$P(w_{1:n}) \approx \prod_{k=1}^n P(w_k|w_{k-1})$$

$$\begin{aligned}P(<S> \text{ i want english food } </S>) \\&= P(\text{i} | <S>)P(\text{want} | \text{i})P(\text{english} | \text{want}) \\&\quad P(\text{food} | \text{english})P(</S> | \text{food}) \\&= .25 \times .33 \times .0011 \times 0.5 \times 0.68 \\&= .000031\end{aligned}$$

A Definition

Language Models assign probabilities to sequences of words.

$$\begin{aligned} P(X_1 \dots X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_{1:2}) \dots P(X_n|X_{1:n-1}) \\ &= \prod_{k=1}^n P(X_k|X_{1:k-1}) \end{aligned}$$

$$P(w_n|w_{n-N+1:n-1}) = \frac{C(w_{n-N+1:n-1} w_n)}{C(w_{n-N+1:n-1})}$$

$$P(w_{1:n}) \approx \prod_{k=1}^n P(w_k|w_{k-1})$$

$$\begin{aligned} P(<S> \text{ i want english food } </S>) \\ &= P(\text{i} | <S>) P(\text{want} | \text{i}) P(\text{english} | \text{want}) \\ &\quad P(\text{food} | \text{english}) P(</S> | \text{food}) \\ &= .25 \times .33 \times .0011 \times 0.5 \times 0.68 \\ &= .000031 \end{aligned}$$

Word Prediction Everywhere

An experiment:

- ▶ Open any chat/messaging app you use frequently
- ▶ Start typing

I wish this lecture was ----

- ▶ What do you get after **was**?
- ▶ The same idea also applies also to full sentences!


Word Prediction Everywhere

An experiment:

- ▶ Open any chat/messaging app you use frequently
- ▶ Start typing

I wish this lecture was ----

- ▶ What do you get after **was**?
- ▶ The same idea also applies also to full sentences!



Word Prediction Everywhere [cont.]

- ▶ Humans do it too!

Please turn your homework ----

Why is automatizing this useful?

- ▶ speech recognition
- ▶ spell-checking/grammatical error correction
- ▶ machine translation
- ▶ maybe even more direct linguistics research ...

This is where Language Models (LMs) enter the picture!

Word Prediction Everywhere [cont.]

- ▶ Humans do it too!

Please turn your homework ----

Why is automatizing this useful?

- ▶ speech recognition
- ▶ spell-checking/grammatical error correction
- ▶ machine translation
- ▶ maybe even more direct linguistics research ...

This is where Language Models (LMs) enter the picture!

Word Prediction Everywhere [cont.]

- ▶ Humans do it too!

Please turn your homework ----

Why is automatizing this useful?

- ▶ speech recognition
- ▶ spell-checking/grammatical error correction
- ▶ machine translation
- ▶ maybe even more direct linguistics research ...

This is where Language Models (LMs) enter the picture!

Tackling Next Word Prediction

A naive solution to next word prediction:

- ▶ We want the **most likely completion(s)**...

Uhm, how do we figure out what is most likely?

- ▶ Idea! Most likely = Most frequent word

- ▶ **Approach:**

- 1 Collect sufficiently large sample of texts (**corpus**)
- 2 For each word (**type**), count how often it occurs in the entire sample (= its number of **tokens**).
- 3 Calculate the **frequency** of the word in the sample:

$$\text{freq}(\text{word}, \text{sample}) = \frac{\text{number of tokens of } \text{word}}{\text{word length of whole } \text{sample}}$$

Tackling Next Word Prediction

A naive solution to next word prediction:

- ▶ We want the **most likely completion(s)**...

Uhm, how do we figure out what is most likely?

- ▶ Idea! Most likely = Most frequent word

- ▶ **Approach:**

- 1 Collect sufficiently large sample of texts (**corpus**)
- 2 For each word (**type**), count how often it occurs in the entire sample (= its number of **tokens**).
- 3 Calculate the **frequency** of the word in the sample:

$$\text{freq}(\text{word}, \text{sample}) = \frac{\text{number of tokens of } \text{word}}{\text{word length of whole } \text{sample}}$$

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered predictions:

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered predictions: be

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered predictions: be, bell

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered predictions: be, bell, bed

Example calculation

Sample: 1000 words long

Words: be, bed, bee, bell

Type	be	bed	bee	bell
Tokens	13	2	0	3

$$\text{freq}(\text{be}) = \frac{13}{1000} = 1.3\%$$

$$\text{freq}(\text{bee}) = \frac{0}{1000} = 0.0\%$$

$$\text{freq}(\text{bed}) = \frac{2}{1000} = 0.2\%$$

$$\text{freq}(\text{bell}) = \frac{3}{1000} = 0.3\%$$

Ordered predictions: be, bell, bed, bee

Tackling Next Word Prediction [cont.]

We want the **most likely completion(s)**...

- ▶ Idea! Most likely = Most frequent word?
- ▶ BUT! Word usage varies by **context**!

Example

	tested	testing	testimony
I have			
I have been			
I have the			

- ▶ The frequency of words is not enough,
we need frequencies of sequences of words \Rightarrow **n-grams**

Tackling Next Word Prediction [cont.]

We want the **most likely completion(s)**...

- ▶ Idea! Most likely = Most frequent word?
- ▶ BUT! Word usage varies by **context**!

Example

	tested	testing	testimony
I have	hi	low	mid
I have been			
I have the			

- ▶ The frequency of words is not enough,
we need frequencies of sequences of words \Rightarrow **n-grams**

Tackling Next Word Prediction [cont.]

We want the **most likely completion(s)**...

- ▶ Idea! Most likely = Most frequent word?
- ▶ BUT! Word usage varies by **context**!

Example

	tested	testing	testimony
I have	hi	low	mid
I have been	hi	hi	low
I have the	low	low	hi

- ▶ The frequency of words is not enough,
we need frequencies of sequences of words \Rightarrow **n-grams**

Tackling Next Word Prediction [cont.]

We want the **most likely completion(s)**...

- ▶ Idea! Most likely = Most frequent word?
- ▶ BUT! Word usage varies by **context**!

Example

	tested	testing	testimony
I have	hi	low	mid
I have been	hi	hi	low
I have the	low	low	hi

- ▶ The frequency of words is not enough,
we need frequencies of sequences of words \Rightarrow **n-grams**

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

N-Grams for Next Word Prediction

Frequencies can be computed and used for n-grams, too.

→ we still need a representative corpus...

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is __

► Sorted completions

► To predict a word w :

- 1 **Needed resources:** corpus
- 2 Compute frequencies for all n-grams
- 3 Look at previous $n - 1$ words
- 4 Find completions that maximize n-gram probability

N-Grams for Next Word Prediction

Frequencies can be computed and used for n-grams, too.

→ we still need a representative corpus...

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is __

► Sorted completions

late

► To predict a word w :

- 1 **Needed resources:** corpus
- 2 Compute frequencies for all n-grams
- 3 Look at previous $n - 1$ words
- 4 Find completions that maximize n-gram probability

N-Grams for Next Word Prediction

Frequencies can be computed and used for n-grams, too.

→ we still need a representative corpus...

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is __

► Sorted completions

late, lovely

► To predict a word w :

- 1 **Needed resources:** corpus
- 2 Compute frequencies for all n-grams
- 3 Look at previous $n - 1$ words
- 4 Find completions that maximize n-gram probability

N-Grams for Next Word Prediction

Frequencies can be computed and used for n-grams, too.

→ we still need a representative corpus...

Example

► Trigram frequencies

bus is late	30%	train is late	15%
bus is lovely	25%	train is lovely	8%
bus is lazy	10%	train is lazy	2%

► Input

I will text you if the train is __

► Sorted completions

late, lovely, lazy

► To predict a word w :

- 1 **Needed resources:** corpus
- 2 Compute frequencies for all n-grams
- 3 Look at previous $n - 1$ words
- 4 Find completions that maximize n-gram probability

The n-Gram Hypothesis (aka Markov Assumption)

The **preceding $n - 1$ words** reliably predict the next word.

$$P(w_n | w_1 w_2 \dots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$
$$P(\text{late} | I \text{ will text you if the train is}) \approx P(\text{late} | \text{train is})$$

- ▶ The n-gram hypothesis is **not quite satisfying**, though.
 - 1 We are not going to see all possible words in all contexts
 - 2 Many dependencies in language are not local

The n-Gram Hypothesis (aka Markov Assumption)

The **preceding $n - 1$ words** reliably predict the next word.

$$P(w_n | w_1 w_2 \dots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$
$$P(\text{late} | I \text{ will text you if the train is}) \approx P(\text{late} | \text{train is})$$

- ▶ The n-gram hypothesis is **not quite satisfying**, though.
 - 1 We are not going to see all possible words in all contexts
 - 2 Many dependencies in language are not local

The n-Gram Hypothesis (aka Markov Assumption)

The **preceding $n - 1$ words** reliably predict the next word.

$$P(w_n | w_1 w_2 \dots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$
$$P(\text{late} | I \text{ will text you if the } \text{train is}) \approx P(\text{late} | \text{train is})$$

- ▶ The n-gram hypothesis is **not quite satisfying**, though.
 - 1 We are not going to see all possible words in all contexts
 - 2 Many dependencies in language are not local

Linguistic Evaluation

The n-Gram Hypothesis (aka Markov Assumption)

The **preceding** $n - 1$ **words** reliably predict the next word.

$$P(w_n | w_1 w_2 \dots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$
$$P(\text{late} | I \text{ will text you if the } \text{train is}) \approx P(\text{late} | \text{train is})$$

- ▶ The n-gram hypothesis is **not quite satisfying**, though.
 - 1 We are not going to see all possible words in all contexts
 - 2 Many dependencies in language are not local

This is where Neural Networks LMs come in handy...



Ok but, who cares?

LMs assign probabilities to sequences of words.

- ▶ **n-Gram LMs**: use local contexts for sequence prediction
- ▶ Spoilers: **Neural LMs**...

And?

- ▶ speech recognition
- ▶ spell-checking/grammatical error correction
- ▶ text generation (think chatbots)
- ▶ machine translation
- ▶ maybe even (less application-oriented) linguistic research ...

Ok but, who cares?

LMs assign probabilities to sequences of words.

- ▶ **n-Gram LMs**: use local contexts for sequence prediction
- ▶ Spoilers: **Neural LMs**...

And?

- ▶ speech recognition
- ▶ spell-checking/grammatical error correction
- ▶ text generation (think chatbots)
- ▶ machine translation
- ▶ maybe even (less application-oriented) linguistic research ...

Ok but, who cares?

LMs assign probabilities to sequences of words.

- ▶ **n-Gram LMs**: use local contexts for sequence prediction
- ▶ Spoilers: **Neural LMs**...

And?

- ▶ speech recognition
- ▶ spell-checking/grammatical error correction
- ▶ text generation (think chatbots)
- ▶ machine translation
- ▶ maybe even (less application-oriented) linguistic research ...

LMs as Tools for Psycholinguistics

Jacobs, De Santo, and Grobol (2023)

Zeugma The architect bit the lime and **the dust**

Literal The architect bit the lime and the apple

- ▶ We can use LMs to generate literal continuations
The architect bit the ----
- ▶ Maze Task


LMs as Tools for Psycholinguistics

Jacobs, De Santo, and Grobol (2023)

Zeugma The architect bit **the lime** and **the dust**

Literal The architect bit **the lime** and **the apple**

- ▶ We can use LMs to generate literal continuations
The architect bit the ----
- ▶ Maze Task (Boyce & Levy, 2021):
Use LMs to generate low probability foils



LMs as Tools for Sociolinguistics

Making “fetch” happen: The influence of social and linguistic context on nonstandard word growth and decline

Ian Stewart and Jacob Eisenstein
School of Interactive Computing
Georgia Institute of Technology



- ▶ Does the social context of a word influence its adoption more than its linguistic context?
- ▶ Use *unique* n-gram counts to measure *dissemination*: the diversity of linguistic contexts in which a word appears
- ▶ How do communities (e.g. $r/x,y,z$) predict word usage? (Lucy & Bamman, 2021)

LMs as Psycholinguistic Subjects

“Wait...Maybe I find **the models** interesting?”

- ▶ Can we use linguistic tests to understand them better?


LMs as Psycholinguistic Subjects

“Wait...Maybe I find **the models** interesting?”

- Can we use linguistic tests to understand them better?

Agreement attraction errors

- (1)
- a. The **key** **is** on the table.
 - b. *The **key** **are** on the table.
 - c. *The **keys** **is** on the table.
 - d. The **keys** **are** on the table.



Assessing the Ability of LSTMs to Learn Syntax-Sensitive Dependencies

Tal Linzen^{1,2} Emmanuel Dupoux¹
LSCP¹ & IJN², CNRS,
EHESS and ENS, PSL Research University
{tal.linzen,
emmanuel.dupoux}@ens.fr

Yoav Goldberg
Computer Science Department
Bar Ilan University
yoav.goldberg@gmail.com

A Final Note: A Word of Caution

- ▶ LMs are sensitive to statistical regularities in language data...
 - ▶ Bias: treating language behavior as ground truth (Bolukbasi et al. 2016)
 - ▶ Exclusion/discrimination: what kind of data is included? (Bender et al. 2019)
 - ▶ Privacy: whose data and how do we get it? (Huang & Paul 2019)
 - ▶ Environmental and financial cost (Strubell et al. 2019)
 - ▶ And more!
- ▶ Reflect on **social impact** while conducting research!

A Final Note: A Word of Caution

- ▶ LMs are sensitive to statistical regularities in language data...
 - ▶ Bias: treating language behavior as ground truth (Bolukbasi et al. 2016)
 - ▶ Exclusion/discrimination: what kind of data is included? (Bender et al. 2019)
 - ▶ Privacy: whose data and how do we get it? (Huang & Paul 2019)
 - ▶ Environmental and financial cost (Strubell et al. 2019)
 - ▶ And more!
- ▶ Reflect on **social impact** while conducting research!

The End (?)



Appendix

LMs assign probabilities to sequences of words.

- ▶ **n-Gram LMs**: use local contexts for sequence prediction.
 - ▶ Struggle to generalize to novel contexts
 - ▶ Struggle with long distance relations (Markov assumption)
- ▶ Spoilers: **Neural LMs**...
 - ▶ ...might help with these issues
 - ▶ Incorporate word similarity based on distributional information
 - ▶ More complex approximation of sentential dependencies

LMs assign probabilities to sequences of words.

- ▶ **n-Gram LMs**: use local contexts for sequence prediction.
 - ▶ Struggle to generalize to novel contexts
 - ▶ Struggle with long distance relations (Markov assumption)
- ▶ Spoilers: **Neural LMs**...
 - ▶ ...might help with these issues
 - ▶ Incorporate word similarity based on distributional information
 - ▶ More complex approximation of sentential dependencies

LMs assign probabilities to sequences of words.

- ▶ **n-Gram LMs**: use local contexts for sequence prediction.
 - ▶ Struggle to generalize to novel contexts
 - ▶ Struggle with long distance relations (Markov assumption)
- ▶ Spoilers: **Neural LMs**...
 - ▶ ...might help with these issues
 - ▶ Incorporate word similarity based on distributional information
 - ▶ More complex approximation of sentential dependencies

Generalizing to Novel Contexts

Imagine our model has seen sequences like:

*I have to make sure that the cat gets fed.
Pearl's parrot gets fed every day.*

Then we want to complete the following:

I forgot to make sure that the dog gets ----

- It would be great if the model could take advantage of the similarity between *dog,cat,parrot* to predict *fed*!

Generalizing to Novel Contexts

Imagine our model has seen sequences like:

*I have to make sure that the cat gets fed.
Pearl's parrot gets fed every day.*

Then we want to complete the following:


I forgot to make sure that the dog gets ----

- It would be great if the model could take advantage of the similarity between *dog, cat, parrot* to predict *fed*!

From Counts to Vector Spaces

The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.

2-Dimensional Vector Space with *dog* and *cat*



	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	0
run	1	2	0	-

Semantic similarity as angle between vectors:

- *bark* more closely related to *dog*
- *run* more closely related to *cat*

Long-distance Dependencies in Language

- ▶ Word choice can be influenced by words that are very far away.

Subject-verb agreement

- ▶ The key to the cabinet **is** on the table.
 - ▶ The keys to the cabinets **are** on the table.
 - ▶ The key to the cabinets **is/are** on the table.
 - ▶ The keys to the cabinet **is/are** on the table.
-
- ▶ Observation: humans get those “wrong” sometimes...
 - ▶ It’s not just about complex “syntactic” dependencies
I spread like strawberries, I climb like peas and beans
I’ve been sucking it in so long, That I’m busting at the seams

Long-distance Dependencies in Language

- ▶ Word choice can be influenced by words that are very far away.

Subject-verb agreement

- ▶ The key to the cabinet **is** on the table.
 - ▶ The keys to the cabinets **are** on the table.
 - ▶ The key to the cabinets **is/are** on the table.
 - ▶ The keys to the cabinet **is/are** on the table.
-
- ▶ Observation: humans get those “wrong” sometimes...
 - ▶ It’s not just about complex “syntactic” dependencies
I spread like strawberries, I climb like peas and beans
I’ve been sucking it in so long, That I’m busting at the seams

Long-distance Dependencies in Language

- ▶ Word choice can be influenced by words that are very far away.


Subject-verb agreement

- ▶ The key to the cabinet **is** on the table.
 - ▶ The keys to the cabinets **are** on the table.
 - ▶ The key to the cabinets **is/are** on the table.
 - ▶ The keys to the cabinet **is/are** on the table.
-
- ▶ Observation: humans get those “wrong” sometimes...
 - ▶ It’s not just about complex “syntactic” dependencies
*I spread like strawberries, I climb like peas and **beans***
*I’ve been sucking it in so long, That I’m **busting** at the **seams***


A Quick Excursus: The Perceptron

The Perceptron: A Mini-Version of a Neural Network


- ▶ **input layer:** neurons that are sensitive to input
- ▶ **output layer:** neurons that represent output values
- ▶ **connections:** weighted links between input and output layer
- ▶ most activated output neuron represents decision




Perceptron Activation for *Hi Dear*



Perceptron Activation for *Hi Dear Emily*




Putting Things Together: Neural LMs




A Bit More on Conditional Probability

- We said we are interested in $P(\text{late}|\text{is})$

$$P(A|B) = \frac{P(A, B)}{P(B)}$$



- E.g. $P(\text{blue}|\blacksquare) = 2/5$



Estimating Bigram Probabilities: MLE

Ok but where do we get probabilities from?

- ▶ One possibility: Counts (Maximum Likelihood Estimate)!
 - ▶ For a unigram:

$$P(w_n) = \frac{\text{count}(w_n)}{\sum_{w \in V} \text{count}(w)}$$

- ▶ MLE of conditional probability for bigrams:

$$P(w_n | w_{n-1}) = \frac{\text{count}(w_n, w_{n-1})}{\text{count}(w_{n-1})}$$

- ▶ Note that the **normalization factor** is different than what we did for pure bigram frequency counts (which gave us an estimate of **joint probability** for each bigram)!

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,
buffalo buffalo,

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

- ▶ **String**

when buffalo buffalo buffalo buffalo buffalo buffalo

- ▶ **Bigram token list**

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,
buffalo buffalo, buffalo buffalo

- ▶ **Bigram counts and frequencies**

Frequencies for n-grams

Frequencies can be computed for n-grams, too.

Example: Calculating Bigram Frequencies

► String

when buffalo buffalo buffalo buffalo buffalo

► Bigram token list

when buffalo, buffalo buffalo, buffalo buffalo, buffalo buffalo,
buffalo buffalo, buffalo buffalo


► Bigram counts and frequencies

- 1 when buffalo: $1 \Rightarrow \frac{1}{6} = 16.7\%$
- 2 buffalo buffalo: $5 \Rightarrow \frac{5}{6} = 83.3\%$

Ok, NOW Some Formulas

$$P(w_n | w_1 w_2 \cdots w_{n-1})$$

$P(\text{late} | \text{I will text you if the train is})$ $P(\text{lazy} | \text{I will text you if the train is})$



► Lots of possible sentences!

► Simplifying assumption:

$$P(w_n | w_1 w_2 \cdots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$

$$P(\text{late} | \text{I will text you if the train is}) \approx P(\text{late} | \text{train is})$$


The n-Gram Hypothesis (aka Markov Assumption)

One can reliably predict the next word based on the **preceding $n - 1$ words**.

Ok, NOW Some Formulas

$$P(w_n | w_1 w_2 \cdots w_{n-1})$$

$P(\text{late} | \text{I will text you if the train is})$ $P(\text{lazy} | \text{I will text you if the train is})$



- ▶ Lots of possible sentences!

- ▶ Simplifying assumption:

$$P(w_n | w_1 w_2 \cdots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$

$$P(\text{late} | \text{I will text you if the train is}) \approx P(\text{late} | \text{train is})$$


The n-Gram Hypothesis (aka Markov Assumption)

One can reliably predict the next word based on the **preceding $n - 1$ words**.

Ok, NOW Some Formulas

$$P(w_n | w_1 w_2 \cdots w_{n-1})$$

$P(\text{late} | \text{I will text you if the train is})$ $P(\text{lazy} | \text{I will text you if the train is})$



- ▶ Lots of possible sentences!

- ▶ Simplifying assumption:

$$P(w_n | w_1 w_2 \cdots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$

$$P(\text{late} | \text{I will text you if the train is}) \approx P(\text{late} | \text{train is})$$


The n-Gram Hypothesis (aka Markov Assumption)

One can reliably predict the next word based on the **preceding $n - 1$ words**.

Ok, NOW Some Formulas

$$P(w_n | w_1 w_2 \cdots w_{n-1})$$

$P(\text{late} | I \text{ will text you if the train is})$ $P(\text{lazy} | I \text{ will text you if the train is})$



- ▶ Lots of possible sentences!
- ▶ Simplifying assumption:

$$P(w_n | w_1 w_2 \cdots w_{n-2} w_{n-1}) \approx P(w_n | w_{n-2} w_{n-1})$$

$$P(\text{late} | I \text{ will text you if the } \text{train is}) \approx P(\text{late} | \text{train is})$$

The n-Gram Hypothesis (aka Markov Assumption)

One can reliably predict the next word based on the **preceding $n - 1$ words**.

An Observation on Frequencies: Zipf's Law

- ▶ Word models care about word frequency.
- ▶ But there is a problem...

Zipf's Law

The frequency of a type is inversely proportional to its rank.



In Plain English

The most frequent word is

- ▶ 2 times as common as the 2nd most frequent word,
- ▶ 3 times as common as the 3rd most frequent word,
- ▶ and so on.

An Observation on Frequencies: Zipf's Law

- ▶ Word models care about word frequency.
- ▶ But there is a problem...

Zipf's Law

The frequency of a type is inversely proportional to its rank.




In Plain English


The most frequent word is

- ▶ 2 times as common as the 2nd most frequent word,
- ▶ 3 times as common as the 3rd most frequent word,
- ▶ and so on.

An Example from...the NBA?




Visualizing Zipf Distributions



Zipf's Law is Everywhere. . .

- ▶ A distribution is probably Zipfian if
 - ▶ there is a **long neck**:
a few types make up the majority of tokens,
 - ▶ there is a **long tail**:
most types only have 1 token (**hapax legomenon**)
- ▶ Surprisingly, Zipf's Law shows up in tons of places:
 - ▶ size of large cities in a country
 - ▶ citations for academic papers
 - ▶ frequencies of last names
 - ▶ frequencies of weekdays in text

... Even in Language!



An Important Consequence of Zipf's Law

- ▶ Texts mostly consist of stop words.
- ▶ Hence it can be difficult to get representative counts for non-stop words.

Sparse Data Problem

- ▶ Most of the data is not informative.
- ▶ You need tons of data to have enough useful data.

An Important Consequence of Zipf's Law

- ▶ Texts mostly consist of stop words.
- ▶ Hence it can be difficult to get representative counts for non-stop words.

Sparse Data Problem

- ▶ Most of the data is not informative.
- ▶ You need tons of data to have enough useful data.

Example

- ▶ Most models require corpora with at least a few million sentences.
- ▶ Really good models (e.g. Google translate) use billions of data points.

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

Defining Larger n-Grams

n-gram a contiguous sequence of n words

n	Name	Example
1	unigram	John
2	bigram	John to
3	trigram	John to be
4	4-gram	John to be in
5	5-gram	John to be in the

Example

String

John and Marie are not Bill and Sue

How long can n-grams be?

- ▶ It is tempting to move to longer and longer n-grams in order to handle long-distance dependencies.
- ▶ But this has **two problems**:
 - data sparsity longer n-grams require too much data
 - storage needs longer n-grams require lots of storage
- ▶ Data sparsity is much more severe than storage needs.

Sparse data: A simple calculation

Words	bigrams	trigrams	5-grams	6-grams
10	100	1000	10,000	100,000
100	10,000	1,000,000	10,000,000,000	1,000,000,000,000
10,000	10^8	10^{12}	10^{20}	10^{24}
25,000	6.3×10^8	1.6×10^{13}	9.7×10^{21}	2.4×10^{26}

Some comparison values

4.3×10^{17} number of seconds since the Big Bang

5×10^{22} number of stars in observable universe

10^{24} milliliters of water in the Earth's oceans

8.8×10^{26} diameter of observable universe, in meters

10^{80} number of atoms in observable universe

- **Conclusion:** with large n , most n -grams are **never encountered** in a corpus \Rightarrow frequency 0

Things get worse: A more realistic estimate

- ▶ The Linux dictionary `american-english-insane` has 650,000 entries.
- ▶ This makes the numbers much worse.
Can you guess how many 5-grams there are then?

Things get worse: A more realistic estimate

- ▶ The Linux dictionary american-english-insane has 650,000 entries.
- ▶ This makes the numbers much worse.
Can you guess how many 5-grams there are then?

$$116 \text{ octillion} \approx 10^{29}$$

Things get worse: A more realistic estimate

- ▶ The Linux dictionary american-english-insane has 650,000 entries.
- ▶ This makes the numbers much worse.
Can you guess how many 5-grams there are then?

$$116 \text{ octillion} \approx 10^{29}$$

10^{29} is larger than the number of shotglasses it takes to drain the Earth's oceans over 2000 times.

Evaluating Language Models: Perplexity

The **perplexity** of a language model is defined as the inverse of the probability of the test set, normalized by the number of tokens (N) in the test set.

$$PP(w_1...w_N) = \sqrt[N]{\frac{1}{P(w_1...w_N)}}$$

A LM with lower perplexity is better because it assigns a higher probability to the unseen test corpus. But note that two LMs can be compared wrt to perplexity iff they use the same vocabulary!

- ▶ Trigram models have lower perplexity than bigram models, etc.

Intrinsic vs. Extrinsic Evaluation

Perplexity tells us which LM assigns a higher probability to unseen text.

This doesn't necessarily tell us which LM is better for a specific task.

Task-based evaluation:

- ▶ Train model A, plug it into your system for performing task T
- ▶ Evaluate performance of system A on task T
- ▶ Train model B, plug it in, evaluate system B on same task T
- ▶ Compare scores of system A and system B on task T.

Originally developed for speech recognition.

How much does the *predicted* sequence of words differ from the *actual* sequence of words in the correct transcript?

$$\text{WER} = \frac{\text{Insertions} + \text{Deletions} + \text{Substitutions}}{\text{Actual words in transcript}}$$

Insertions: “eat lunch” → “eat **a** lunch”

Deletions: “see **a** movie” → “see movie”

Substitutions: “drink **ice** tea” → “drink **nice** tea”

¹slide adapted from J. Hockenmaier